

Practical 1

Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset.

Importing libraries and the dataset

```
In [1]: #Importing the pandas for data processing and numpy for numerical computing
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: # Importing the Boston Housing dataset from the sklearn
from sklearn.datasets import fetch_openml

# Load the Boston housing dataset from OpenML
boston = fetch_openml(data_id=531)
```

C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\sklearn\datasets_openml.py:932: FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.

```
warn(
```

```
In [3]: #Converting the data into pandas dataframe
data = pd.DataFrame(boston.data)
```

First look at the dataset

```
In [4]: #First Look at the data
data.head()
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33

```
In [5]: #Adding the feature names to the dataframe
data.columns = boston.feature_names
```

```
In [6]: #Adding the target variable to the dataset
data['PRICE'] = boston.target
```

```
In [7]: #Looking at the data with names and target variable
data.head()
```

```
Out[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2

```
In [8]: #Shape of the data
print(data.shape)
```

```
(506, 14)
```

```
In [9]: #Checking the null values in the dataset
data.isnull().sum()
```

```
Out[9]: CRIM      0
ZN          0
INDUS      0
CHAS       0
NOX        0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT      0
PRICE      0
dtype: int64
```

No null values in the dataset, no missing value treatment needed

```
In [10]: #Checking the statistics of the data
data.describe()
```

```
Out[10]:
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	TAX
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.554695	6.284634	68.574901	3.795043	408.237154
std	8.601545	23.322453	6.860353	0.115878	0.702617	28.148861	2.105710	168.537116
min	0.006320	0.000000	0.460000	0.385000	3.561000	2.900000	1.129600	187.000000
25%	0.082045	0.000000	5.190000	0.449000	5.885500	45.025000	2.100175	279.000000
50%	0.256510	0.000000	9.690000	0.538000	6.208500	77.500000	3.207450	330.000000
75%	3.677083	12.500000	18.100000	0.624000	6.623500	94.075000	5.188425	666.000000
max	88.976200	100.000000	27.740000	0.871000	8.780000	100.000000	12.126500	711.000000

This is sometimes very useful, for example if you look at the CRIM the max is 88.97 and 75% of the value is below 3.677083 and mean is 3.613524 so it means the max values is actually an outlier or there are outliers present in the column


```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 506 entries, 0 to 505  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype    
---  -  
0   CRIM        506 non-null    float64  
1   ZN          506 non-null    float64  
2   INDUS       506 non-null    float64  
3   CHAS        506 non-null    category  
4   NOX         506 non-null    float64  
5   RM          506 non-null    float64  
6   AGE         506 non-null    float64  
7   DIS         506 non-null    float64  
8   RAD         506 non-null    category  
9   TAX         506 non-null    float64  
10  PTRATIO     506 non-null    float64  
11  B           506 non-null    float64  
12  LSTAT       506 non-null    float64  
13  PRICE       506 non-null    float64  
dtypes: category(2), float64(12)  
memory usage: 49.0 KB
```

Visualisation

```
In [12]: #checking the distribution of the target variable  
import seaborn as sns  
sns.displot(data['PRICE'], kde=True)
```


```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x216a12168f0>
```



The distribution seems normal, has not be the data normal we would have perform log transformation or took to square root of the data to make the data normal. Normal distribution is need for the machine learning for better predictiblity of the model

```
In [13]: #Distribution using box plot
sns.boxplot(data.PRICE)
```

```
Out[13]: <Axes: >
```



Checking the correlation of the independent feature with the dependent feature

Correlation is a statistical technique that can show whether and how strongly pairs of variables are related. An intelligent correlation analysis can lead to a greater understanding of your data

```
In [14]: #checking Correlation of the data
correlation = data.corr()
correlation.loc['PRICE']
```


C:\Users\Yashwardhan Deshmukh\AppData\Local\Temp\ipykernel_11992\1690761071.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation = data.corr()
```

```
Out[14]: CRIM      -0.388305
          ZN        0.360445
          INDUS    -0.483725
          NOX      -0.427321
          RM       0.695360
          AGE      -0.376955
          DIS       0.249929
          TAX      -0.468536
          PTRATIO  -0.507787
          B        0.333461
          LSTAT    -0.737663
          PRICE     1.000000
          Name: PRICE, dtype: float64
```


```
In [15]: # plotting the heatmap
import matplotlib.pyplot as plt
fig, axes = plt.subplots(figsize=(15,12))
sns.heatmap(correlation, square = True, annot = True)
```

Out[15]: <Axes: >



By looking at the correlation plot LSTAT is negatively correlated with -0.74 and RM is positively correlated to the price and PTRATIO is correlated negatively with -0.51

```
In [16]: # Checking the scatter plot with the most correlated features
plt.figure(figsize = (20,5))
features = ['LSTAT', 'RM', 'PTRATIO']
for i, col in enumerate(features):
    plt.subplot(1, len(features) , i+1)
    x = data[col]
    y = data.PRICE
    plt.scatter(x, y, marker='o')
    plt.title("Variation in House prices")
    plt.xlabel(col)
    plt.ylabel('"House prices in $1000"')
```



Splitting the dependent feature and independent feature

```
In [17]: #X = data[['LSTAT', 'RM', 'PTRATIO']]
X = data.iloc[:, :-1]
y= data.PRICE
```

Splitting the data for Model Validation

```
In [18]: # Splitting the data into train and test for building the model
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2, random_state=42)
```

Building the Model

```
In [19]: #Linear Regression
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
```

```
In [20]: #Fitting the model
regressor.fit(X_train,y_train)
```

```
Out[20]:
LinearRegression
LinearRegression()
```

Model Evaluation

```
In [21]: # Convert X_test to a NumPy array
X_test = np.array(X_test)
#Prediction on the test dataset
y_pred = regressor.predict(X_test)
```

C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(

```
In [22]: # Predicting RMSE the Test set results
from sklearn.metrics import mean_squared_error
rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
print(rmse)
```

5.041784121402046

```
In [23]: from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print(r2)
```

0.7263451459702515

Neural Networks

```
In [24]: !pip install tensorflow
!pip install keras
```

mukh\anaconda3\lib\site-packages (from tensorflow==2.15.0->tensorflow-intel==2.15.0->tensorflow) (2.23.4)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (0.7.2)
Requirement already satisfied: google-auth-oauthlib<2,>=0.5 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (1.1.0)
Requirement already satisfied: requests<3,>=2.21.0 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (2.28.1)
Requirement already satisfied: markdown<=2.6.8 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (3.4.1)
Requirement already satisfied: werkzeug<=1.0.1 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (2.2.2)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in c:\users\yashwardhan deshmukh\anaconda3\lib\site-packages (from google-auth<3,>=1.6.3->tensorboard<2.16,>=2.15->tensorflow-intel==2.15.0->tensorflow) (5.3.2)

```
In [25]: #Creating the neural network model
import keras
from keras.layers import Dense, Activation, Dropout
from keras.models import Sequential
```

```
model = Sequential()

model.add(Dense(128,activation = 'relu',input_dim =13))
model.add(Dense(64,activation = 'relu'))
model.add(Dense(32,activation = 'relu'))
model.add(Dense(16,activation = 'relu'))
model.add(Dense(1))
model.compile(optimizer = 'adam',loss = 'mean_squared_error')
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

```
In [26]: #Scaling the dataset
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(

```
In [27]: results=model.fit(X_train, y_train, epochs = 100)
```

```
Epoch 91/100
13/13 [=====] - 0s 3ms/step - loss: 3.9828
Epoch 92/100
13/13 [=====] - 0s 2ms/step - loss: 4.0298
Epoch 93/100
13/13 [=====] - 0s 2ms/step - loss: 3.8867
Epoch 94/100
13/13 [=====] - 0s 2ms/step - loss: 3.9619
Epoch 95/100
13/13 [=====] - 0s 2ms/step - loss: 3.7390
Epoch 96/100
13/13 [=====] - 0s 2ms/step - loss: 3.8308
Epoch 97/100
13/13 [=====] - 0s 3ms/step - loss: 3.9280
Epoch 98/100
13/13 [=====] - 0s 3ms/step - loss: 3.6752
Epoch 99/100
13/13 [=====] - 0s 3ms/step - loss: 3.6492
Epoch 100/100
13/13 [=====] - 0s 3ms/step - loss: 3.6492
```


Evaluation of the model

```
In [28]: y_pred = model.predict(X_test)
```

4/4 [=====] - 0s 2ms/step


```
In [29]: from sklearn.metrics import r2_score  
r2 = r2_score(y_test, y_pred)  
print(r2)
```

0.8848322436428823

```
In [30]: # Predicting RMSE the Test set results  
from sklearn.metrics import mean_squared_error  
rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))  
print(rmse)
```

3.270755808161393

```
In [31]: plt.plot(results.history['loss'])  
plt.title('Model Loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.show()
```




```
In [32]: from keras.layers import Dropout
from keras import regularizers

model_3 = Sequential([
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01), input_shape=(784,)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1, activation='sigmoid', kernel_regularizer=regularizers.l2(0.01)),
])
```

```
In [33]: model_3.compile(optimizer='adam',
                        loss='binary_crossentropy',
                        metrics=['accuracy'])
hist_3 = model_3.fit(X_train, y_train,
                    batch_size=32, epochs=100,
                    validation_data=(X_test, y_test))

13/13 [=====] - 1s 46ms/step - loss: -3292891833171968.0
000 - accuracy: 0.0000e+00 - val_loss: -3737344175767552.0000 - val_accuracy: 0.0
000e+00
Epoch 78/100
13/13 [=====] - 1s 48ms/step - loss: -3512145081794560.0
000 - accuracy: 0.0000e+00 - val_loss: -3980877948256256.0000 - val_accuracy: 0.0
000e+00
Epoch 79/100
13/13 [=====] - 1s 48ms/step - loss: -3720911899328512.0
000 - accuracy: 0.0000e+00 - val_loss: -4232192859308032.0000 - val_accuracy: 0.0
000e+00
Epoch 80/100
13/13 [=====] - 1s 48ms/step - loss: -3960872259026944.0
000 - accuracy: 0.0000e+00 - val_loss: -4501129484304384.0000 - val_accuracy: 0.0
000e+00
Epoch 81/100
13/13 [=====] - 1s 49ms/step - loss: -4202862796079104.0
000 - accuracy: 0.0000e+00 - val_loss: -4787134040899584.0000 - val_accuracy: 0.0
000e+00
Epoch 82/100
```

```
In [34]: plt.plot(hist_3.history['loss'])
plt.plot(hist_3.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



Practical 2

Classification using Deep neural network: Binary classification using Deep Neural Networks

Example: Classify movie reviews into "positive" reviews and "negative" reviews, just based on the text content of the reviews. Use IMDB dataset

Import the required libraries

```
In [1]: !pip install git+https://github.com/tensorflow/tensorflow.git@v2.15.0
!pip install keras
```

```
Collecting git+https://github.com/tensorflow/tensorflow.git@v2.15.0
  Cloning https://github.com/tensorflow/tensorflow.git (https://github.com/tensorflow/tensorflow.git) (to revision v2.15.0) to c:\users\yashwardhan deshmunh\appdata\local\temp\pip-req-build-75a22au7
```

```
Running command git clone --filter=blob:none --quiet https://github.com/tensorflow/tensorflow.git (https://github.com/tensorflow/tensorflow.git) 'C:\Users\Yashwardhan Deshmukh\AppData\Local\Temp\pip-req-build-75a22au7'
error: 27943 bytes of body are still expected
fetch-pack: unexpected disconnect while reading sideband packet
fatal: early EOF
fatal: fetch-pack: invalid index-pack output
error: subprocess-exited-with-error
```

```
git clone --filter=blob:none --quiet https://github.com/tensorflow/tensorflow.git (https://github.com/tensorflow/tensorflow.git) 'C:\Users\Yashwardhan Deshmukh\AppData\Local\Temp\pip-req-build-75a22au7' did not run successfully.
exit code: 128
```

See above for output.

```
note: This error originates from a subprocess, and is likely not a problem with pip.
error: subprocess-exited-with-error
```

```
git clone --filter=blob:none --quiet https://github.com/tensorflow/tensorflow.git (https://github.com/tensorflow/tensorflow.git) 'C:\Users\Yashwardhan Deshmukh\AppData\Local\Temp\pip-req-build-75a22au7' did not run successfully.
exit code: 128
```

See above for output.

note: This error originates from a subprocess, and is likely not a problem with pip.

Requirement already satisfied: keras in c:\users\yashwardhan deshmunh\anaconda3\lib\site-packages (2.15.0)

```
In [2]: import pandas as pd
import numpy as np
import keras
import tensorflow as tf
from matplotlib import pyplot as plt
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

Read the Data

```
In [3]: df_train=pd.read_csv('Train.csv')
df_val=pd.read_csv('Valid.csv')
df_train.head()
```

```
Out[3]:
```

	text	label
0	I grew up (b. 1965) watching and loving the Th...	0
1	When I put this movie in my DVD player, and sa...	0
2	Why do people who do not know what a particula...	0
3	Even though I have great interest in Biblical ...	0
4	Im a die hard Dads Army fan and nothing will e...	1

```
In [4]: df_val.head()
```

```
Out[4]:
```

	text	label
0	It's been about 14 years since Sharon Stone aw...	0
1	someone needed to make a car payment... this i...	0
2	The Guidelines state that a comment must conta...	0
3	This movie is a muddled mish-mash of clichés f...	0
4	Before Stan Laurel became the smaller half of ...	0

```
In [5]: X_train=df_train['text'].values
Y_train=df_train['label'].values
```

```
In [6]: X_val=df_val['text'].values
Y_val=df_val['label'].values
```

```
In [7]: (X_train.shape,Y_train.shape),(X_val.shape,Y_val.shape)
```

```
Out[7]: (((40000,), (40000,)), ((5000,), (5000,)))
```

Analyse the Data

```
In [8]: df_train.iloc[:,1].describe()
```

```
Out[8]: count    40000.000000
mean         0.499525
std          0.500006
min          0.000000
25%          0.000000
50%          0.000000
75%          1.000000
max          1.000000
Name: label, dtype: float64
```

```
In [9]: df_val.iloc[:,1].describe()
```

```
Out[9]: count    5000.000000
      mean      0.502800
      std      0.500042
      min      0.000000
      25%      0.000000
      50%      1.000000
      75%      1.000000
      max      1.000000
      Name: label, dtype: float64
```

```
In [10]: X_val_len=[len(str(i).split()) for i in X_val]
      df1=pd.DataFrame(X_val_len,columns=['len'])
      df1.describe()
```

```
Out[10]:
```

	len
count	5000.00000
mean	228.93260
std	169.33721
min	10.00000
25%	126.00000
50%	171.00000
75%	274.00000
max	1601.00000


```
In [11]: X_train_len=[len(str(i).split()) for i in X_train]
      df=pd.DataFrame(X_train_len,columns=['len'])
      df.describe()
```

```
Out[11]:
```

	len
count	40000.000000
mean	231.339250
std	171.194123
min	4.000000
25%	126.000000
50%	173.000000
75%	282.000000
max	2470.000000


```
In [12]: X_train_len=[len(str(i).split()) for i in X_train]
plt.hist(X_train_len)
```

```
Out[12]: (array([2.8097e+04, 8.6960e+03, 2.2520e+03, 8.0100e+02, 1.3800e+02,
                7.0000e+00, 3.0000e+00, 3.0000e+00, 1.0000e+00, 2.0000e+00]),
array([ 4. , 250.6, 497.2, 743.8, 990.4, 1237. , 1483.6, 1730.2,
        1976.8, 2223.4, 2470. ]),
<BarContainer object of 10 artists>)
```



```
In [13]: X_val_len=[len(str(i).split()) for i in X_val]
plt.hist(X_val_len)
```

```
Out[13]: (array([2.470e+03, 1.602e+03, 5.350e+02, 1.970e+02, 1.180e+02, 5.100e+01,
                2.300e+01, 1.000e+00, 2.000e+00, 1.000e+00]),
          array([ 10. , 169.1, 328.2, 487.3, 646.4, 805.5, 964.6, 1123.7,
                1282.8, 1441.9, 1601. ]),
          <BarContainer object of 10 artists>)
```



Setting the parameters

```
In [14]: from tensorflow.keras.datasets import imdb
# Load the IMDB movie review dataset
(X_train, y_train), (X_val, y_val) = imdb.load_data()

# Load word index to convert integer sequences back to text
word_index = imdb.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

# Retrieve the actual text data from integer sequences
X_train_text = [' '.join([reverse_word_index.get(i - 3, '') for i in sequence]) for sequence in X_train]
X_val_text = [' '.join([reverse_word_index.get(i - 3, '') for i in sequence]) for sequence in X_val]

# Define tokenizer and preprocess the sequences
vocab_size = 30000 #went for an average vocab size
embedding_dimension = 64 #high dimensions would result in finding better parameters for
max_length = 120 #used a maximum length of 120 words
turnc = 'post' #preprocessing step for pad_sequences
oov_tok = '<OOV>' #oov stands for out of vocabulary tokens
```


Tokenizing and converting the data into Sequences

```
In [15]: tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(X_train_text)

X_train_seq = tokenizer.texts_to_sequences(X_train_text)
X_train_padded = pad_sequences(X_train_seq, maxlen=max_length, padding='post', truncating='post')

X_val_seq = tokenizer.texts_to_sequences(X_val_text)
X_val_padded = pad_sequences(X_val_seq, maxlen=max_length, padding='post', truncating='post')
```

```
In [16]: X_train_padded.shape, X_val_padded.shape
```

```
Out[16]: ((25000, 120), (25000, 120))
```

The Model

```
In [17]: from tensorflow.keras.layers import LSTM, Bidirectional, Embedding, Dense, SpatialDropout1D
from tensorflow.keras.models import Sequential

# Define the model
model = Sequential()
embedding_layer = Embedding(input_dim=vocab_size, output_dim=embedding_dimension, input_length=sequence_length)
model.add(embedding_layer)
model.add(SpatialDropout1D(0.4))
model.add(Bidirectional(LSTM(120, activation='tanh', return_sequences=True)))
model.add(Dropout(0.3))
model.add(Bidirectional(LSTM(120, activation='tanh', return_sequences=False)))
model.add(Dropout(0.2))
model.add(Dense(300, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])

# Print model summary
print(model.summary())
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\optimizers_init_.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 64)	1920000
spatial_dropout1d (Spatial Dropout1D)	(None, 120, 64)	0
bidirectional (Bidirectional)	(None, 120, 240)	177600
dropout (Dropout)	(None, 120, 240)	0
bidirectional_1 (Bidirectional)	(None, 240)	346560
dropout_1 (Dropout)	(None, 240)	0
dense (Dense)	(None, 300)	72300
dropout_2 (Dropout)	(None, 300)	0
dense_1 (Dense)	(None, 1)	301

=====
Total params: 2516761 (9.60 MB)
Trainable params: 2516761 (9.60 MB)
Non-trainable params: 0 (0.00 Byte)

None

```
In [18]: hist=model.fit(X_train_padded,y_train,epochs=7,batch_size=16,validation_data=(X_val_p
```

Epoch 1/7

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

1563/1563 [=====] - 794s 497ms/step - loss: 0.5323 - accuracy: 0.7200 - val_loss: 0.3863 - val_accuracy: 0.8290

Epoch 2/7

1563/1563 [=====] - 691s 442ms/step - loss: 0.3555 - accuracy: 0.8502 - val_loss: 0.3885 - val_accuracy: 0.8173

Epoch 3/7

1563/1563 [=====] - 650s 416ms/step - loss: 0.3143 - accuracy: 0.8745 - val_loss: 0.3784 - val_accuracy: 0.8286

Epoch 4/7

1563/1563 [=====] - 679s 435ms/step - loss: 0.2795 - accuracy: 0.8894 - val_loss: 0.6437 - val_accuracy: 0.7637

Epoch 5/7

1563/1563 [=====] - 955s 611ms/step - loss: 0.2483 - accuracy: 0.9058 - val_loss: 0.3735 - val_accuracy: 0.8371

Epoch 6/7

1563/1563 [=====] - 919s 588ms/step - loss: 0.2270 - accuracy: 0.9177 - val_loss: 0.4115 - val_accuracy: 0.8309

Epoch 7/7

1563/1563 [=====] - 622s 398ms/step - loss: 0.2004 - accuracy: 0.9272 - val_loss: 0.4894 - val_accuracy: 0.8236

```
In [19]: hist=model.fit(X_train_padded,y_train,epochs=2,batch_size=16,validation_data=(X_val_p
```

Epoch 1/2


1563/1563 [=====] - 510s 327ms/step - loss: 0.1773 - accuracy: 0.9362 - val_loss: 0.5264 - val_accuracy: 0.8256

Epoch 2/2


1563/1563 [=====] - 548s 351ms/step - loss: 0.1552 - accuracy: 0.9460 - val_loss: 0.5695 - val_accuracy: 0.8239

This Plot is for the last two Epochs

```
In [20]: plt.plot(hist.history['accuracy'],c='b')
plt.plot(hist.history['val_accuracy'],c='r')
plt.show()
```



```
In [21]: plt.plot(hist.history['loss'],c='b')
plt.plot(hist.history['val_loss'],c='r')
plt.show()
```



Reading the Test Data

```
In [22]: df_test=pd.read_csv('Test.csv')
df_test.head()
```

```
Out[22]:
```

	text	label
0	I always wrote this series off as being a comp...	0
1	1st watched 12/7/2002 - 3 out of 10(Dir-Steve ...	0
2	This movie was so poorly written and directed ...	0
3	The most interesting thing about Miryang (Secr...	1
4	when i first read about "berlin am meer" i did...	0

```
In [23]: X_test=df_test['text'].values
Y_test=df_test['label'].values
```

Converting into Sequential Data

```
In [24]: X_test_seq=tokenizer.texts_to_sequences(X_test)
```

```
In [25]: X_test_padded=pad_sequences(X_test_seq,maxlen=max_length,padding='post',truncating='tu...
X_test_padded[0]
```

```
Out[25]: array([[ 11,   208,  1040,   12,   199,   123,   15,   110,    4,
                  598,  7221,  3338,   86,  1239,  5390,   14,   573,    9,
                   10,    3,  2697,   19,   93,   29,  249,    4,  1576,
                13809,  3847,   101,    4,   426,  5744,   498,  1052,   11,
                1703,   13,    2,  2951,   14,   30,    2,   94,   21,
                   2,   83,   497,    5,    2,   671,   818,  148,   11,
                   98,   26,   41,  1838,   54,    3,  2530,   636,    2,
                  671,    6,   77,    2,  2951,   40,   58,    6,    2,
                  246,    6,  469,    2,  1306,   19,   93,   136,   22,
                   41,   77,   54,    3,  1133,   636,    2,   702,    6,
                  104,   246,    9,   158,  1109,  4764,    6,   13,   11,
                  299,   36,   11,   870,    6,   41,  3250,  2704,   21,
                   2,  5734,    3,   191,   842,  1947,   67,    9,  1132,
                   16,   70,   49])
```

```
In [26]: X_test_padded.shape
```

```
Out[26]: (5000, 120)
```

```
In [27]: model.evaluate(X_test_padded,Y_test)
```

```
157/157 [=====] - 22s 140ms/step - loss: 0.3230 - accuracy: 0.8952
```

```
Out[27]: [0.32298383116722107, 0.8952000141143799]
```

Check for your own Reviews

```
In [34]: def Check(x):
          turnc = 'post' # Define the truncating method, e.g., 'post' or 'pre'
          test_case1 = [x]
          test_case = tokenizer.texts_to_sequences(test_case1)
          max_sequence_length = 120 # Define the maximum sequence length expected by the model
          test_case_padded = pad_sequences(test_case, maxlen=max_sequence_length, padding='post')
          predict_x = model.predict(test_case_padded)
          print(predict_x)
          if predict_x >= 0.5:
              print("Positive")
          else:
              print("Negative")
```

```
In [37]: test_review=str(input("Enter the review : "))
          Check(test_review)
```

Enter the review : I learned about this movie couple of weeks ago through a facebook page. The storyline caught my attention and decided to watch it.and am glad i did. It is dark,dramatic and a bit of humor.I had to mention humor because it wasn't like those comedy scenes that you find in other Bollywood films.It made sense and that went with the story.'Shaitan' is realistic in every sense. Movie has five main characters and the whole story revolves around the 'events' that they did.It shows the life of 5 young friends who is living life like hell.They drink,they shop lift,they fool around,pretty much everything that makes you feel alive.But a bit too much adrenal in rush followed by some unwise decisions turns their life around.

1/1 [=====] - 0s 47ms/step

[[0.99900836]]

Positive

I just checked for one random imdb review

```
In [38]: test_review="You will get A to Z all details of this scam, i may be wrong but due to it"
          Check(test_review)
```

1/1 [=====] - 0s 51ms/step

[[0.07417779]]

Negative

Practical 3

Convolutional neural network (CNN): Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories.


Import the required libraries

```
In [ ]: #!pip install git+https://github.com/tensorflow/tensorflow.git@v2.15.0  
        #!pip install keras
```

Fashion MNIST Classification

Fashion-MNIST is a dataset of Zalando's article images—consisting of a **training set of 60,000** examples and a **test set of 10,000 examples**. Each example is a **28x28 grayscale** image, associated with a label from **10 classes**. We intend Fashion-MNIST to serve as a direct drop-in **replacement for the original MNIST** dataset for benchmarking machine learning algorithms. It shares the same image size and structure of training and testing splits.

Here's an example how the data looks (each class takes three-rows):



Step # 1 - Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sbn
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dropout, Dense, Flatten
from keras.optimizers import Adam
from keras.callbacks import TensorBoard
from keras.utils import to_categorical
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Step # 2 - Load Data

```
In [2]: fashion_train_df = pd.read_csv('dataset/fashion-mnist_train.csv', sep=',')
fashion_test_df = pd.read_csv('dataset//fashion-mnist_test.csv', sep=',')
```

Now that we have loaded the datasets, lets check some parameters about the datasets.

```
In [3]: fashion_train_df.shape    # Shape of the dataset
```

```
Out[3]: (60000, 785)
```

```
In [4]: fashion_train_df.columns  # Name of the columns of the DataSet.
```

```
Out[4]: Index(['label', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5', 'pixel6',
              'pixel7', 'pixel8', 'pixel9',
              ...,
              'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779', 'pixel780',
              'pixel781', 'pixel782', 'pixel783', 'pixel784'],
              dtype='object', length=785)
```

```
In [5]: print(set(fashion_train_df['label']))
```

```
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

```
In [6]: print([fashion_train_df.drop(labels='label', axis=1).min(axis=1).min(),
              fashion_train_df.drop(labels='label', axis=1).max(axis=1).max()])
```

```
[0, 255]
```

So we have 0 to 255 which is the color values for grayscale. 0 being white and 255 being black.

Now lets check some of the rows in tabular format


```
In [7]: fashion_train_df.head()
```

Out[7]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777
0	2	0	0	0	0	0	0	0	0	0	...	0	0	0
1	9	0	0	0	0	0	0	0	0	0	...	0	0	0
2	6	0	0	0	0	0	0	0	5	0	...	0	0	0
3	0	0	0	0	1	2	0	0	0	0	...	3	0	0
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0

5 rows × 785 columns



So every other thing of the test dataset are going to be the same as the train dataset except the shape.

```
In [8]: fashion_test_df.shape
```

Out[8]: (10000, 785)

So here we have 10000 images instead of 60000 as in the train dataset.

Let's check first few rows.

```
In [9]: fashion_test_df.head()
```

Out[9]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777
0	0	0	0	0	0	0	0	0	9	8	...	103	87	5
1	1	0	0	0	0	0	0	0	0	0	...	34	0	0
2	2	0	0	0	0	0	0	14	53	99	...	0	0	0
3	2	0	0	0	0	0	0	0	0	0	...	137	126	14
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0

5 rows × 785 columns



Step # 3 - Visualization

Now that we have loaded the data and also got somewhat acquainted with it let's visualize the actual images. We are going to use **Matplotlib** library for this.


```
In [10]: # Convert the dataframe to numpy array
training = np.asarray(fashion_train_df, dtype='float32')

# Lets show multiple images in a 15x15 grid
height = 10
width = 10

fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(17,17))
axes = axes.ravel() # this flattens the 15x15 matrix into 225
n_train = len(training)

for i in range(0, height*width):
    index = np.random.randint(0, n_train)
    axes[i].imshow(training[index, 1:].reshape(28,28))
    axes[i].set_title(int(training[index, 0]), fontsize=8)
    axes[i].axis('off')

plt.subplots_adjust(hspace=0.5)
```



Step # 4 - Preprocess Data

Visualized the images. So now we can start preparing for creating our model. But before that we need to preprocess our data so that we can fit our model easily. Lets do that first.

Since we are dealing with image data and our task is to recognize and classify images our model should be a Convolutional Neural Network. For that our images should have atleast 3 dimensions (**height x width x color_channels**). But our images are flattened in one dimension, **784 pixel (28x28x1)** values per row. So we need to reshape the data into its original format

```
In [11]: # convert to numpy arrays and reshape
training = np.asarray(fashion_train_df, dtype='float32')
X_train = training[:, 1:].reshape([-1,28,28,1])
X_train = X_train/255 # Normalizing the data
y_train = training[:, 0]

testing = np.asarray(fashion_test_df, dtype='float32')
X_test = testing[:, 1:].reshape([-1,28,28,1])
X_test = X_test/255 # Normalizing the data
y_test = testing[:, 0]
```

Also we need to have three different sets of data for **training**, **validatin** and **testing**. We already have different sets for training and testing. So we are going to split the training dataset further into two sets and will use one set of training and the other for validation.

```
In [12]: # Split the training set into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, ra
```

```
In [13]: # Lets check the shape of all three datasets
print(X_train.shape, X_val.shape, X_test.shape)
print(y_train.shape, y_val.shape, y_test.shape)

(48000, 28, 28, 1) (12000, 28, 28, 1) (10000, 28, 28, 1)
(48000,) (12000,) (10000,)
```

Step # 5 - Create and Train the Model

Create the model

```
In [14]: cnn_model = Sequential()
cnn_model.add(Conv2D(filters=64, kernel_size=(3,3), input_shape=(28,28,1), activation:
cnn_model.add(MaxPooling2D(pool_size = (2,2)))
cnn_model.add(Dropout(rate=0.3))
cnn_model.add(Flatten())
cnn_model.add(Dense(units=32, activation='relu'))
cnn_model.add(Dense(units=10, activation='sigmoid'))
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

compile the model

```
In [15]: cnn_model.compile(optimizer=Adam(learning_rate=0.001), loss='sparse_categorical_crossentropy', metrics=['accuracy'])
cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d (MaxPooling2D)	(None, 13, 13, 64)	0
dropout (Dropout)	(None, 13, 13, 64)	0
flatten (Flatten)	(None, 10816)	0
dense (Dense)	(None, 32)	346144
dense_1 (Dense)	(None, 10)	330

=====
Total params: 347114 (1.32 MB)
Trainable params: 347114 (1.32 MB)
Non-trainable params: 0 (0.00 Byte)
=====

Train the model

```
In [16]: cnn_model.fit(x=X_train, y=y_train, batch_size=512, epochs=50, validation_data=(X_val, y_val))
```

94/94 [=====] - 47s 502ms/step - loss: 0.1344 - accuracy: 0.9523 - val_loss: 0.2502 - val_accuracy: 0.9178
Epoch 45/50
94/94 [=====] - 44s 467ms/step - loss: 0.1326 - accuracy: 0.9524 - val_loss: 0.2528 - val_accuracy: 0.9173
Epoch 46/50
94/94 [=====] - 41s 434ms/step - loss: 0.1341 - accuracy: 0.9507 - val_loss: 0.2559 - val_accuracy: 0.9142
Epoch 47/50
94/94 [=====] - 44s 468ms/step - loss: 0.1280 - accuracy: 0.9540 - val_loss: 0.2514 - val_accuracy: 0.9165
Epoch 48/50
94/94 [=====] - 45s 483ms/step - loss: 0.1267 - accuracy: 0.9542 - val_loss: 0.2551 - val_accuracy: 0.9158
Epoch 49/50
94/94 [=====] - 43s 462ms/step - loss: 0.1251 - accuracy: 0.9543 - val_loss: 0.2518 - val_accuracy: 0.9164
Epoch 50/50
94/94 [=====] - 43s 457ms/step - loss: 0.1253 - accuracy: 0.9547 - val_loss: 0.2481 - val_accuracy: 0.9200

Step # 5 - Evaluate the Model

Get the accuracy of the model

```
In [17]: eval_result = cnn_model.evaluate(X_test, y_test)
print("Accuracy : {:.3f}".format(eval_result[1]))
```

313/313 [=====] - 5s 14ms/step - loss: 0.2446 - accuracy: 0.9210
Accuracy : 0.921


Visualize the model's predictions

```
In [18]: predict_x=cnn_model.predict(X_test)
classes_x=np.argmax(predict_x,axis=1)
```

313/313 [=====] - 6s 12ms/step

```
In [19]: height = 10
width = 10


fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(20,20))
axes = axes.ravel()
for i in range(0, height*width):
    index = np.random.randint(len(classes_x))
    axes[i].imshow(X_test[index].reshape((28,28)))
    axes[i].set_title("True Class : {:.0f}\nPrediction : {:.0f}".format(y_test[index],
    axes[i].axis('off')
plt.subplots_adjust(hspace=0.9, wspace=0.5)
```



Plot Confusin Matrix

```
In [20]: cm = confusion_matrix(y_test, classes_x)
plt.figure(figsize=(10,5))
sbn.heatmap(cm, annot=True)
```

Out[20]: <Axes: >



Classification Report

```
In [21]: num_classes = 10
class_names = ["class {}".format(i) for i in range(num_classes)]
cr = classification_report(y_test, classes_x, target_names=class_names)
print(cr)
```

	precision	recall	f1-score	support
class 0	0.84	0.90	0.87	1000
class 1	0.98	0.99	0.99	1000
class 2	0.91	0.84	0.87	1000
class 3	0.93	0.93	0.93	1000
class 4	0.88	0.90	0.89	1000
class 5	0.99	0.97	0.98	1000
class 6	0.78	0.77	0.78	1000
class 7	0.94	0.98	0.96	1000
class 8	0.99	0.98	0.98	1000
class 9	0.97	0.95	0.96	1000
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

Practical 4

Recurrent neural network (RNN): Use the Google stock prices dataset and design a timeseriesanalysis and prediction system using RNN.

1. Import library

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

2. Data processing

2.0 import the data

```
In [2]: dataset_train = pd.read_csv('Google_Stock_Price_Train.csv')
```

```
In [3]: dataset_train.head()
```

Out[3]:


	Date	Open	High	Low	Close	Volume
0	01/03/2012	325.25	332.83	324.97	663.59	7,380,500
1	01/04/2012	331.27	333.87	329.08	666.45	5,749,400
2	01/05/2012	329.83	330.75	326.89	657.21	6,590,300
3	01/06/2012	328.34	328.77	323.68	648.24	5,405,900
4	01/09/2012	322.04	322.29	309.46	620.76	11,688,800

```
In [4]: #keras only takes numpy array
training_set = dataset_train.iloc[:, 1: 2].values
```

```
In [5]: training_set.shape
```

Out[5]: (1509, 1)

```
In [6]: plt.figure(figsize=(18, 8))
plt.plot(dataset_train['Open'])
plt.title("Google Stock Open Prices")
plt.xlabel("Time (oldest -> latest)")
plt.ylabel("Stock Open Price")
plt.show()
```



2.1 Feature scaling

```
In [7]: import os
if os.path.exists('config.py'):
    print(1)
else:
    print(0)
```

0

```
In [8]: sc = MinMaxScaler(feature_range = (0, 1))
#fit: get min/max of train data
training_set_scaled = sc.fit_transform(training_set)
```

2.2 Data structure creation

- taking the reference of past 60 days of data to predict the future stock price.
- It is observed that taking 60 days of past data gives us best results.
- In this data set 60 days of data means 3 months of data.
- Every month as 20 days of Stock price.
- X train will have data of 60 days prior to our date and y train will have data of one day after our date

```
In [9]: ## 60 timesteps and 1 output
X_train = []
y_train = []
for i in range(60, len(training_set_scaled)):
    X_train.append(training_set_scaled[i-60: i, 0])
    y_train.append(training_set_scaled[i, 0])

X_train, y_train = np.array(X_train), np.array(y_train)
```



```
In [10]: X_train.shape
```

```
Out[10]: (1449, 60)
```

```
In [11]: y_train.shape
```

```
Out[11]: (1449,)
```

2.3 Data reshaping

```
In [12]: X_train = np.reshape(X_train, newshape =  
                                (X_train.shape[0], X_train.shape[1], 1))
```

1. Number of stock prices - 1449
2. Number of time steps - 60
3. Number of Indicator - 1

```
In [13]: X_train.shape
```

```
Out[13]: (1449, 60, 1)
```

3. Create & Fit Model

3.1 Create model

```
In [14]: regressor = Sequential()  
#add 1st Lstm Layer  
regressor.add(LSTM(units = 50, return_sequences = True, input_shape = (X_train.shape[0], X_train.shape[1], 1)))  
regressor.add(Dropout(rate = 0.2))  
  
##add 2nd Lstm Layer: 50 neurons  
regressor.add(LSTM(units = 50, return_sequences = True))  
regressor.add(Dropout(rate = 0.2))  
  
##add 3rd Lstm Layer  
regressor.add(LSTM(units = 50, return_sequences = True))  
regressor.add(Dropout(rate = 0.2))  
  
##add 4th Lstm Layer  
regressor.add(LSTM(units = 50, return_sequences = False))  
regressor.add(Dropout(rate = 0.2))  
  
##add output Layer  
regressor.add(Dense(units = 1))
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

```
In [15]: regressor.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 60, 50)	10400
dropout (Dropout)	(None, 60, 50)	0
lstm_1 (LSTM)	(None, 60, 50)	20200
dropout_1 (Dropout)	(None, 60, 50)	0
lstm_2 (LSTM)	(None, 60, 50)	20200
dropout_2 (Dropout)	(None, 60, 50)	0
lstm_3 (LSTM)	(None, 50)	20200
dropout_3 (Dropout)	(None, 50)	0
dense (Dense)	(None, 1)	51

=====
Total params: 71051 (277.54 KB)
Trainable params: 71051 (277.54 KB)
Non-trainable params: 0 (0.00 Byte)

```
In [16]: regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

WARNING:tensorflow:From C:\Users\Yashwardhan Deshmukh\anaconda3\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

3.2 Model fit

```
In [17]: regressor.fit(x = X_train, y = y_train, batch_size = 32, epochs = 100)
```

```
Epoch 92/100
46/46 [=====] - 13s 274ms/step - loss: 9.1957e-04
Epoch 93/100
46/46 [=====] - 13s 282ms/step - loss: 9.3358e-04
Epoch 94/100
46/46 [=====] - 13s 294ms/step - loss: 0.0010
Epoch 95/100
46/46 [=====] - 14s 310ms/step - loss: 9.5691e-04
Epoch 96/100
46/46 [=====] - 14s 300ms/step - loss: 8.3259e-04
Epoch 97/100
46/46 [=====] - 12s 261ms/step - loss: 9.6896e-04
Epoch 98/100
46/46 [=====] - 11s 234ms/step - loss: 8.9036e-04
Epoch 99/100
46/46 [=====] - 11s 236ms/step - loss: 9.6050e-04
Epoch 100/100
46/46 [=====] - 11s 237ms/step - loss: 9.2182e-04
```

3.3 Model evaluation

3.3.1 Read and convert

```
In [18]: dataset_test = pd.read_csv('Google_Stock_Price_Test.csv')
```

```
In [19]: dataset_test.head()
```

Out[19]:

	Date	Open	High	Low	Close	Volume
0	02/01/2018	1048.339966	1066.939941	1045.229980	1065.000000	1237600
1	03/01/2018	1064.310059	1086.290039	1063.209961	1082.479980	1430200
2	04/01/2018	1088.000000	1093.569946	1084.001953	1086.400024	1004600
3	05/01/2018	1094.000000	1104.250000	1092.000000	1102.229980	1279100
4	08/01/2018	1102.229980	1111.270020	1101.619995	1106.939941	1047600

```
In [20]: #keras only takes numpy array  
real_stock_price = dataset_test.iloc[:, 1: 2].values  
real_stock_price.shape
```

Out[20]: (125, 1)

3.3.2 Concat and convert

```
In [21]: #vertical concat use 0, horizontal uses 1  
dataset_total = pd.concat((dataset_train['Open'], dataset_test['Open']),  
                           axis = 0)  
  
##use .values to make numpy array  
inputs = dataset_total[len(dataset_total) - len(dataset_test) - 60:].values
```

3.3.3 Reshape and scale

```
In [22]: #reshape data to only have 1 col  
inputs = inputs.reshape(-1, 1)  
  
#scale input  
inputs = sc.transform(inputs)
```

```
In [23]: len(inputs)
```

Out[23]: 185

3.3.4 Create test data strucutre

```
In [24]: X_test = []  
for i in range(60, len(inputs)):  
    X_test.append(inputs[i-60:i, 0])  
X_test = np.array(X_test)  
#add dimension of indicator  
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

```
In [25]: X_test.shape
```

Out[25]: (125, 60, 1)

3.3.5 Model prediction

```
In [26]: predicted_stock_price = regressor.predict(X_test)
```

4/4 [=====] - 26s 92ms/step

```
In [27]: #inverse the scaled value  
predicted_stock_price = sc.inverse_transform(predicted_stock_price)
```

3.3.6 Result visualization

```
In [28]: ##visualize the prediction and real price  
plt.plot(real_stock_price, color = 'red', label = 'Real price')  
plt.plot(predicted_stock_price, color = 'blue', label = 'Predicted price')  
  
plt.title('Google price prediction')  
plt.xlabel('Time')  
plt.ylabel('Price')  
plt.legend()  
plt.show()
```

